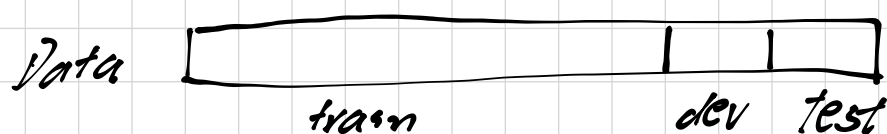


Improve deep neural network - week 1

1. Set up ML Application

(1) Train / Dev / Test set



By Data 99% / 0.5% / 0.5%

enough to judge algorithm

train / test set should match \rightarrow come from same distribution

(2) Bias / Variance

Train set error:

Dev set error:

1%
11%

15%
16%

15%
30%

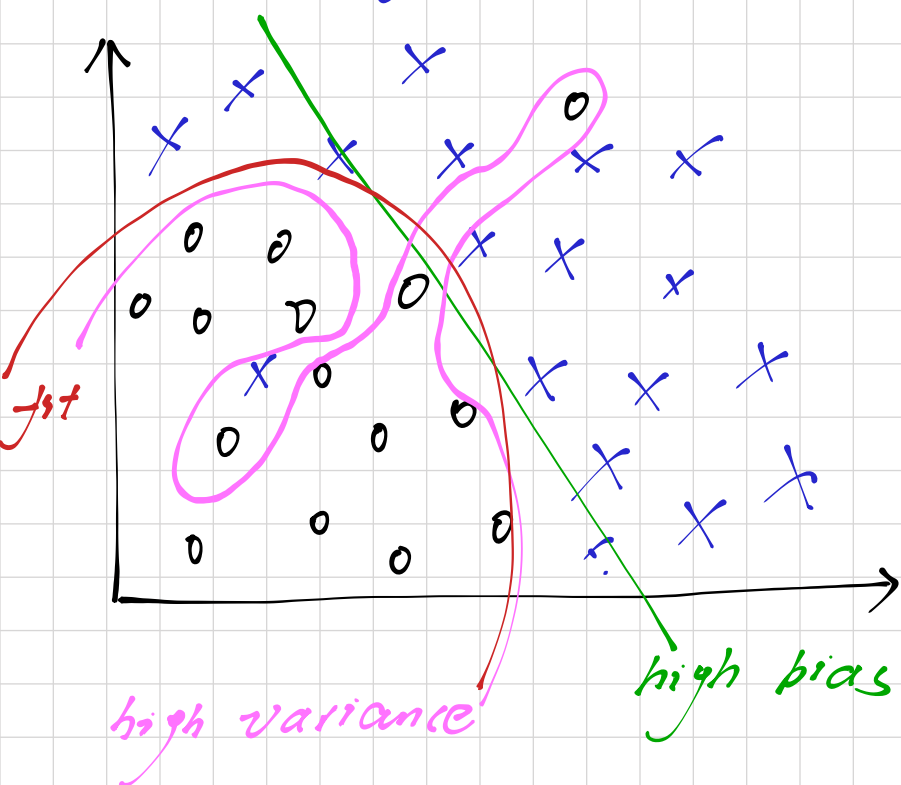
0.5%
1%

\rightarrow low bias & low variance

high variance

high bias

high variance & high bias



(3) Basic recipe

High bias?
(training data performance)

\rightarrow Bigger network
train longer

High variance?
(dev set performance)

\rightarrow More data
Regularization

\downarrow N
done

2. Regularizing NN

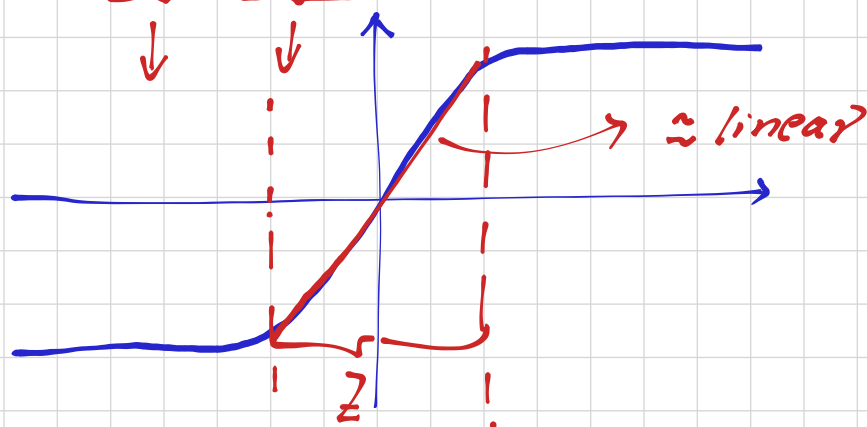
(1) Regularization

$$J(w^{[1]}, b^{[1]}, \dots, w^{[L]}, b^{[L]}) = \frac{1}{m} \sum_{i=1}^n \mathcal{L}(\hat{y}^{(i)}, y^{(i)}) + \frac{\lambda}{2m} \sum_{l=1}^L \|w^{[l]}\|_F^2$$

↓ prevent overfitting

$$\|w^{[L]}\|_F^2 = \sum_{i=1}^{n^{[L-1]}} \sum_{j=1}^{n^{[L]}} (w_{ij}^{[L]})^2$$

$$\lambda \uparrow \rightarrow w^{[L]} \downarrow \rightarrow z^{[L]} = w^{[L]} a^{[L-1]} + b^{[L]}$$



∴ complex → linear?

↓ overfitting

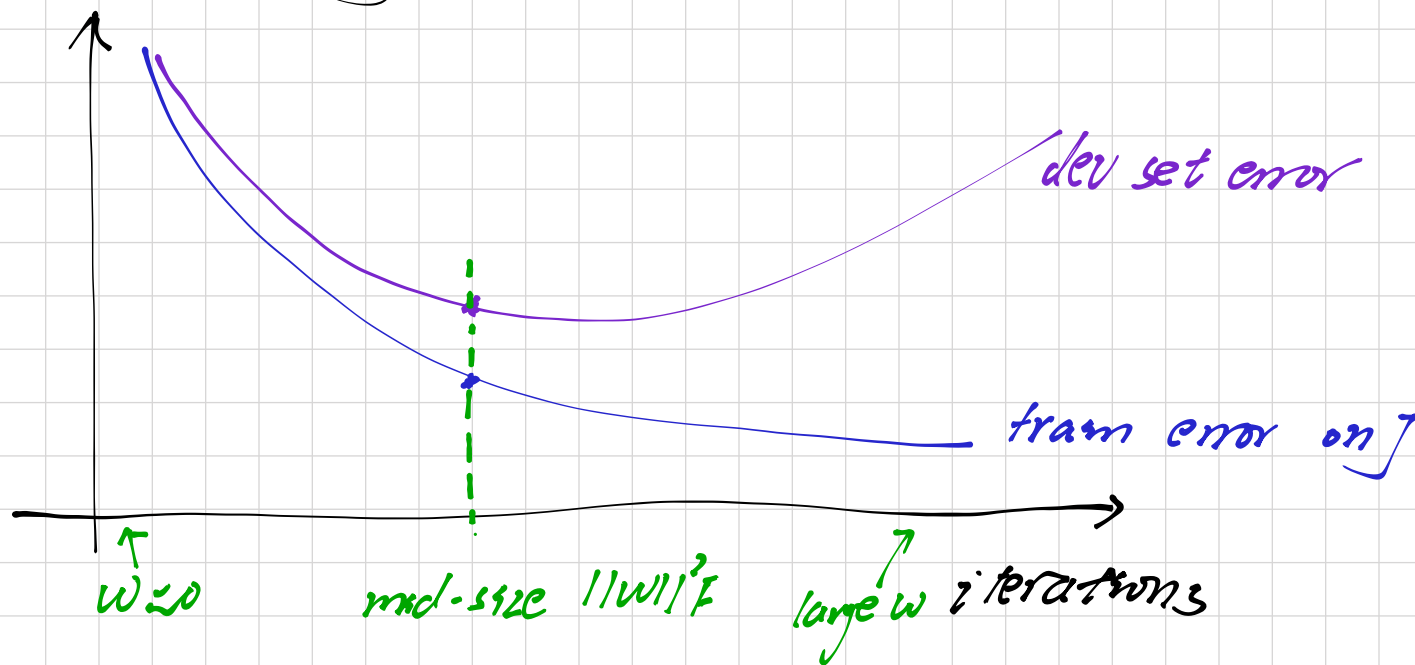
(2) Dropout

① can't rely on any one feature, so have to spread out weights

shrink weights

② choose different keep-prob (depend on the complexity of each layer)

(3) Early stopping



3. Gradient descent check

for each i :

$$\alpha \text{ Approximate}[i] = \frac{J(\theta_1, \theta_2, \dots, \theta_i + \epsilon, \dots) - J(\theta_1, \theta_2, \dots, \theta_i - \epsilon, \dots)}{2\epsilon}$$

$$\approx d\theta[i] = \frac{\partial J}{\partial \theta_i}$$

check $\frac{\|d\theta_{\text{approximate}} - d\theta\|_2}{\|d\theta_{\text{approximate}}\|_2 + \|d\theta\|_2} \approx 10^{-7} \rightarrow \text{great}$
 10^{-5}
 $10^{-3} \rightarrow \text{worry}$

① don't use in training — only to debug

② if algorithm fails grad check, look at components to try to identify bug

③ Remember regularization

④ Don't work with dropout

⑤ Run at random initialization; perhaps again after some training